

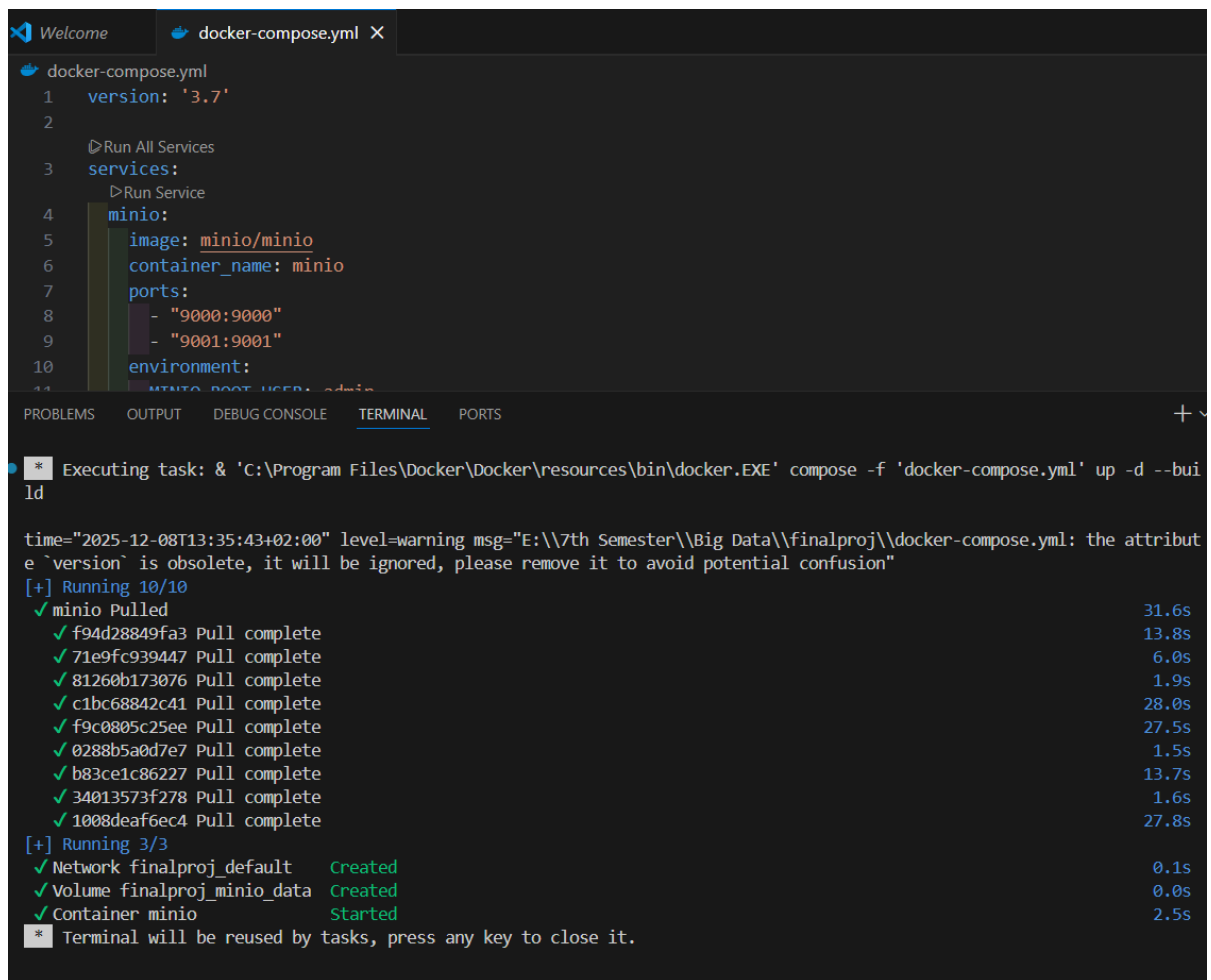
Big Data Project Report

Name	ID
Younna El Sayed Saad	22010297
Hamass Nagieb Mortada	22010331
Malak Ali Ahmed Saed	22010265
Moiad Ishag Abdalla Abdelmolla	22012059
Mohamedahmed Abdelhadi Abdelgadir	22012150
Habiba Mohamed Bassiouny	20221462320
Mohammed Abdallah Abdelhafith	20221460214

Phase 1: Infrastructure & Data Ingestion (Bronze Layer)

1. MinIO Setup:

- Created a Docker Compose file (docker-compose.yml) to start the MinIO server locally.
- Configured MinIO with default access keys.



The screenshot shows a code editor with a file named `docker-compose.yml` and a terminal window below it. The `docker-compose.yml` file contains the following configuration:

```
1 version: '3.7'
2
3 services:
4   minio:
5     image: minio/minio
6     container_name: minio
7     ports:
8       - "9000:9000"
9       - "9001:9001"
10    environment:
```

The terminal window shows the output of the command `docker-compose up -d`. It indicates that the `minio` service was pulled successfully and is now running. The output also shows that the `finalproj_default` network and `finalproj_minio_data` volume were created.

```
time="2025-12-08T13:35:43+02:00" level=warning msg="E:\\7th Semester\\Big Data\\finalproj\\docker-compose.yml: the attribute 'version' is obsolete, it will be ignored, please remove it to avoid potential confusion"
[+] Running 10/10
✓ minio Pulled
  ✓ f94d28849fa3 Pull complete
  ✓ 71e9fc939447 Pull complete
  ✓ 81260b173076 Pull complete
  ✓ c1bc68842c41 Pull complete
  ✓ f9c0805c25ee Pull complete
  ✓ 0288b5a0d7e7 Pull complete
  ✓ b83ce1c86227 Pull complete
  ✓ 34013573f278 Pull complete
  ✓ 1008deaf6ec4 Pull complete
[+] Running 3/3
✓ Network finalproj_default Created
✓ Volume finalproj_minio_data Created
✓ Container minio Started
* Terminal will be reused by tasks, press any key to close it.
```

Containers [Give feedback](#)

Container CPU usage
0.18% / 800% (8 CPUs available)

Container memory usage
85.79MB / 3.66GB

[Show charts](#)




Only show running containers

☐	Name	Container ID	Image	Port(s)	CPU (%)	Actions
☐	finalproj	-	-	-	0.18%	
☐	minio	2d510c72714b	minio/minio	9000:9000 Show all ports (2)	0.18%	

2. Bucket Creation:

Created three buckets in MinIO:


- bronze → stores raw datasets
- silver → will store cleaned data
- gold → will store results and reports



3. Data Upload:

Uploaded the raw synthetic datasets into the **bronze** bucket:

- weather_raw.csv
- traffic_raw.csv



Phase 2: Data Cleaning

Objective:

The goal of this phase was to transform the raw, messy data stored in the Bronze layer into clean, structured, and high-quality datasets stored in the Silver layer using Python and Pandas. The cleaning process involved handling missing values, removing duplicates, correcting data types, and treating outliers.

1. Weather Dataset Cleaning

The following transformation rules were applied to the `weather_raw.csv` dataset:

- **IDs (`weather_id`):**
 - Converted to numeric type and removed exact row duplicates.
 - Handled missing or duplicated IDs by generating new unique sequential IDs starting from the maximum existing ID to ensure data integrity¹.
- **Dates (`date_time`):**
 - Parsed mixed date formats (e.g., "08/05/2024 02AM", ISO format) into a standard datetime object.
 - Removed invalid dates and rows with future dates (beyond 2025).
- **City (`city`):**
 - Filled missing values with "London" as the standard city for this project.
- **Season (`season`):**
 - Standardized text to lowercase.
 - Imputed missing or non-standard values by deriving the correct season from the `date_time` month.
- **Numeric Variables (Outliers Handling):**
 - **Temperature:** Filtered to keep values within a realistic range of -5°C to 40°C.
 - **Humidity:** Restricted values to the valid percentage range (0-100%).
 - **Wind Speed:** Capped at 200 km/h; values exceeding this were removed as outliers.
 - **Visibility:** Filtered to keep values between 50m and 10,000m.
 - **Rainfall (`rain_mm`):** Imputed missing values using the **Median** to avoid skewing data with extreme storms.

-
- **Categorical Variables:**
 - **Weather Condition:** Standardized values to lowercase (e.g., 'rain', 'snow', 'clear') and replaced non-standard or missing values with 'unknown'.

2. Traffic Dataset Cleaning

The following transformation rules were applied to the traffic_raw.csv dataset:

- **IDs (traffic_id):**
 - Similar to the weather dataset, duplicates were removed, and conflicts/nulls were resolved by assigning new unique sequential IDs.
- **Dates (date_time):**
 - Standardized formats and removed future/invalid dates.
- **Location (city & area):**
 - Filled missing city with "London".
 - Filled missing area values with 'unknown' and standardized text to lowercase to ensure consistency.
- **Numeric Variables (Cleaning & Imputation):**
 - **Average Speed:** Converted negative values to absolute values, removed outliers (> 150 km/h), and imputed missing values using the **Mean**.
 - **Vehicle Count:** Converted to absolute values, capped at 5,000 to remove extreme outliers, and imputed missing values using the **Median**.
 - **Accident Count:** Converted to absolute values, filled missing records with 0 (assuming no accident occurred), and removed extreme values (> 10).
 - **Visibility:** Filtered valid range (50-10,000m) and imputed missing values with the **Mean**.
- **Categorical Variables:**
 - **Congestion Level:** Standardized to 'low', 'medium', 'high'. Non-standard or missing values were replaced with 'unknown'.
 - **Road Condition:** Standardized to 'dry', 'wet', 'snowy', 'damaged'. Missing or invalid entries were replaced with 'unknown'.

3. Data Quality Summary

After applying the cleaning pipeline, the datasets were saved as Parquet files in the MinIO **Silver** bucket.

Metric	Weather Dataset	Traffic Dataset
Rows Before Cleaning	5,000	5,000
Rows After Cleaning	3,024	4,151
Actions Taken	Outliers removed, Seasons imputed, formatting fixed.	Negatives fixed, Nulls imputed, Categories standardized.



☐	Name	Last Modified	Size
☐	 traffic_cleaned.parquet	Today, 20:02	1471 KiB
☐	 weather_cleaned.parquet	Today, 20:02	152.4 KiB

Phase 3 Data Transformation & Gold Layer

1. Introduction

Phase 3 focuses on transforming the cleaned data from the Silver Layer and preparing high-quality, analytics-ready datasets for the Gold Layer.

This step is essential to convert the raw cleaned data into structured, aggregated, and feature-enriched datasets suitable for reporting and visualization.

2. Objectives of Phase 3

- Read cleaned Parquet files from the **Silver Layer**
- Apply essential transformations (feature engineering + aggregation)
- Generate Gold Layer datasets
- Upload the final outputs to the **gold bucket** in MinIO

3. Input Datasets (from Silver Layer)

The following cleaned Parquet files were used:

- weather_cleaned.parquet
- traffic_cleaned.parquet

These datasets were produced in Phase 2 after removing errors, missing values, and invalid data.

4. Transformations Performed

4.1 Feature Engineering

To support analytics, additional time-based features were created:

```
[16]: weather['date_time'] = pd.to_datetime(weather['date_time'])
      traffic['date_time'] = pd.to_datetime(traffic['date_time'])

      weather['day'] = weather['date_time'].dt.date
      weather['hour'] = weather['date_time'].dt.hour

      traffic['day'] = traffic['date_time'].dt.date
      traffic['hour'] = traffic['date_time'].dt.hour
```

Purpose:

- Easier daily and hourly analysis
- Support time-series charts
- Improve aggregation and trend detection

4.2 Daily Aggregation (Summary Statistics)

- A daily summary dataset was created to represent average values for each day:


```
[18]: weather_daily = weather.groupby("day").mean(numeric_only=True)
      traffic_daily = traffic.groupby("day").mean(numeric_only=True)
```

Purpose:

- Produce a simplified dataset for dashboards
- Identify trends such as weather changes and traffic patterns
- Reduce granularity for faster insights

5. Output Datasets (Gold Layer)

- Four final datasets were generated and saved as Parquet:
- **weather_gold.parquet**
- Feature-engineered weather dataset
- **traffic_gold.parquet**
- Feature-engineered traffic dataset
- **weather_daily.parquet**
- Daily aggregated weather dataset
- **traffic_daily.parquet**
- Daily aggregated traffic dataset
- These were uploaded to:
- **MinIO → gold bucket**



Name	Last Modified	Size
traffic_daily.parquet	Today, 22:33	19.8 KiB
traffic_gold.parquet	Today, 22:33	157.1 KiB
weather_daily.parquet	Today, 22:33	24.7 KiB
weather_gold.parquet	Today, 22:33	160.4 KiB

6. Hadoop HDFS Integration

As part of Phase 3, we also integrated Hadoop HDFS as an additional storage layer to ensure that the Gold Layer datasets are available in a scalable, distributed environment.

This step supports future big-data processing and aligns with the system architecture used in the project.


6.1 Preparing HDFS for Gold Data

We configured a Hadoop cluster running on Docker, consisting of one NameNode and one DataNode.

After verifying that the cluster was healthy, we created two directories in HDFS to organize the data:

- /weather
- /traffic

These directories were used to store the final output of the Gold Layer.



6.2 Uploading Gold Datasets to HDFS

Once the transformation phase was completed, we uploaded all four Gold Layer Parquet files to HDFS:

- weather_gold.parquet
- weather_daily.parquet
- traffic_gold.parquet
- traffic_daily.parquet

This ensures that the project data is stored not only in MinIO, but also in a distributed file system suitable for large-scale processing.

Browse Directory

/weather Go!   

Show 25 entries Search:

☐	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
☐	-rw-r--r--	root	supergroup	24.69 KB	Dec 11 19:12	3	128 MB	weather_daily.parquet	
☐	-rw-r--r--	root	supergroup	160.38 KB	Dec 11 19:09	3	128 MB	weather_gold.parquet	

Showing 1 to 2 of 2 entries

Previous 1 Next

Hadoop, 2019.

Browse Directory

/traffic Go!   

Show 25 entries Search:


☐	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
☐	-rw-r--r--	root	supergroup	19.84 KB	Dec 11 19:16	3	128 MB	traffic_daily.parquet	
☐	-rw-r--r--	root	supergroup	157.09 KB	Dec 11 19:15	3	128 MB	traffic_gold.parquet	

Showing 1 to 2 of 2 entries

Previous 1 Next

Hadoop, 2019.

• Phase4



```
Copy_of_Untitled6.ipynb X
C:\Users> ahmed > OneDrive > Desktop > phase 4 > Copy_of_Untitled6.ipynb > import os
Generate + Code + Markdown Run All Clear All Outputs Outline ... Select Kernel

weather['date_time'] = pd.to_datetime(weather['date_time'])
traffic['date_time'] = pd.to_datetime(traffic['date_time'])

weather = weather.dropna(subset=['date_time','city'])
traffic = traffic.dropna(subset=['date_time','city'])

weather['city'] = weather['city'].str.strip().str.lower()
traffic['city'] = traffic['city'].str.strip().str.lower()

weather.rename(columns={"datetime": "date_time", inplace=True})
traffic.rename(columns={"datetime": "date_time", inplace=True})

merged = pd.merge(
    weather,
    traffic,
    on=['date_time','city'],
    how='inner'
)

print(weather.shape)
print(traffic.shape)
print(merged.shape)
```

1. Project Overview

The purpose of this project is to analyze how weather conditions affect urban traffic behavior in the city of London.

A modern data lake architecture was designed and implemented to store, clean, process, and analyze large-scale weather and traffic datasets.

The final goal is to support better traffic planning and management decisions under different weather conditions.

2. Technology Stack

The following tools and technologies were used in this project:

- **Programming Language:** Python
- **Infrastructure:** Docker
- **Object Storage:** MinIO
 - Bronze Layer (Raw Data)
 - Silver Layer (Cleaned Data)
 - Gold Layer (Final Results)
- **Distributed File System:** HDFS
- **Data Formats:**
 - CSV (Raw Data)
 - Parquet (Cleaned and Processed Data)

- **Analysis Techniques:**
 - Monte Carlo Simulation
 - Factor Analysis
-

3. Data Lake Architecture

The project follows a multi-layer data lake architecture:

1. **Bronze Layer:** Stores raw and unprocessed CSV datasets.
 2. **Silver Layer:** Stores cleaned and structured datasets in Parquet format.
 3. **HDFS Layer:** Used as an additional distributed storage system for cleaned data.
 4. **Gold Layer:** Stores final analytical outputs and reports.
-

4. Phase 4 – Dataset Merging

Objective

The objective of Phase 4 is to create a unified analytical dataset by merging the cleaned weather and traffic datasets.

This merged dataset is required for advanced analysis in later phases.

Input Datasets

- Cleaned Weather Dataset (weather_cleaned.parquet)
 - Cleaned Traffic Dataset (traffic_cleaned.parquet)
- Both datasets were obtained from the **Silver Layer** after completing the data cleaning phase.
-

Merging Methodology

- The datasets were merged using an **INNER JOIN**.
- The join was performed on the following common keys:
 - date_time
 - city
- Before merging:
 - Date and time columns were converted to a unified datetime format.
 - City names were standardized (lowercase and trimmed).
 - Records with missing merge keys were removed.

This approach ensures that only records with matching weather and traffic data for the same city and time are included.

Output

- The result of the merge is a single unified dataset:
 - **merged_data.parquet**
 - The merged dataset was stored in the **Silver/Gold Layer**.
 - This dataset is ready for use in:
 - Monte Carlo Simulation
 - Factor Analysis
-

5. Conclusion

Phase 4 successfully produced a clean and unified analytical dataset by merging weather and traffic data.

This dataset provides a strong foundation for predictive modeling and statistical analysis in subsequent phases of the project.

6. Next Steps

- Apply Monte Carlo Simulation to estimate traffic congestion and accident risks.
- Perform Factor Analysis to identify the most influential weather factors affecting traffic behavior.
- Generate final insights and recommendations for urban traffic management.

Phase 5: Traffic Risk Prediction using Monte Carlo Simulation

The primary objective of this phase was to move beyond simple historical analysis and predict traffic behavior under hypothetical and extreme weather conditions. We aimed to quantify the probability of Traffic Congestion and High Accident Risk under specific scenarios such as heavy rain, extreme temperatures, and poor visibility.

Methodology

To achieve this, we employed a **Monte Carlo Simulation** approach. The process involved the following key steps:

1. **Data Integration:** Merging traffic and weather datasets to create a unified view of historical conditions.
2. **Model Training:** Developing Machine Learning models (Logistic Regression) to understand the mathematical relationship between weather variables (temperature, rain, wind, etc.) and traffic outcomes.
3. **Simulation:** Generating thousands of synthetic data points representing specific weather scenarios (e.g., "What if it rains 100mm?").
4. **Risk Assessment:** Using the trained models to predict outcomes for these synthetic scenarios to derive probability distributions.

Data Preparation

We began by importing the cleaned datasets from Phase 2 (`traffic_cleaned.parquet` and `weather_cleaned.parquet`).

Step 1: Loading the Data

```
traffic = pd.read_parquet('../Phase2/data_sets_silver/traffic_cleaned.parquet')
[2] ✓ 0.0s

weather = pd.read_parquet('../Phase2/data_sets_silver/weather_cleaned.parquet')
[3] ✓ 0.0s

traffic.describe()
[4] ✓ 0.0s
...
weather.describe()
[5] ✓ 0.0s
...
```

	traffic_id	date_time	vehicle_count	avg_speed_kmh	accident_count	visibility_m
count	4151.000000	4151	4151.000000	4151.000000	4151.000000	4151.000000
mean	14555.132498	2024-06-29 15:13:03.526860800	2503.085762	60.910974	4.386413	6198.059986
min	9001.000000	2024-01-01 04:53:00	1.000000	3.028226	0.000000	50.000000
25%	12641.000000	2024-03-28 17:39:00	1246.500000	31.396957	2.000000	2609.500000
50%	15351.000000	2024-06-28 01:09:00	2505.000000	60.995291	4.000000	5197.000000
75%	16604.500000	2024-09-30 19:45:30	3757.500000	89.214658	7.000000	7672.500000
max	17852.000000	2024-12-30 21:00:00	4999.000000	119.961300	9.000000	50000.000000
std	2525.251769	NaN	1450.806527	33.494657	2.941690	7659.755604

	weather_id	date_time	temperature_c	humidity	rain_mm	wind_speed_kmh	visibility_m
count	3024.000000	3024	3024.000000	3024.000000	3024.000000	3024.000000	3024.000000
mean	10157.293320	2024-06-30 01:24:53.6111111168	20.237532	57.969577	27.290784	40.326014	5754.593585
min	5001.000000	2024-01-01 04:09:00	5.007249	10.000000	0.027336	0.007863	50.000000
25%	8651.000000	2024-03-26 20:45:00	12.799367	38.000000	13.108750	19.658138	2494.250000
50%	10713.000000	2024-06-30 15:45:00	20.267637	57.500000	25.380248	40.959270	5021.000000
75%	11950.250000	2024-10-01 09:49:15	27.756274	78.000000	37.544195	60.173435	7588.000000
max	13176.000000	2024-12-30 23:50:00	34.994298	99.000000	120.000000	79.988524	50000.000000
std	2236.930686	NaN	8.646245	23.841918	20.368702	23.200822	6562.021312

Step 2: The Merged Dataset

```
merged
✓ 0.0s
```

	traffic_id	date_time	city	area	vehicle_count	avg_speed_kmh	accident_count	congestion_level	road_condition	visibility_m_x	weather_id	season	temp_c
0	9371	2024-04-16 00:00:00	London	chelsea	767	49.255082	2	low	snowy	2609	11884	autumn	10.0
1	11972	2024-07-09 11:00:00	London	southwark	1131	72.694131	0	unknown	damaged	9326	11095	winter	10.0
2	15023	2024-04-21 04:00:00	London	unknown	2231	86.181583	3	low	damaged	1584	8731	summer	10.0
3	14255	2024-05-14 14:00:00	London	kensington	1711	44.056048	3	medium	dry	6826	7574	winter	10.0
4	15027	2024-12-14 01:00:00	London	unknown	4930	107.788422	4	high	unknown	3437	12904	spring	10.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...
148	14748	2024-12-01 10:00:00	London	camden	1918	21.160514	9	medium	unknown	8826	10969	winter	10.0
149	11413	2024-07-03 16:00:00	London	southwark	3441	86.477826	0	medium	unknown	5108	10995	winter	10.0
150	11413	2024-07-03 16:00:00	London	southwark	3441	86.477826	0	medium	unknown	5108	11234	spring	10.0
151	17775	2024-06-10 01:00:00	London	chelsea	2089	36.809758	9	unknown	damaged	8130	12285	spring	10.0
152	14797	2024-11-01 16:00:00	London	camden	1313	102.056977	4	medium	dry	4916	10063	summer	10.0

Exploratory Data Analysis (EDA)

Before running simulations, we analyzed the distribution of the original data to understand the baseline conditions for variables like temperature, rainfall, and traffic volume.



Simulation Results

We trained Logistic Regression models and ran simulations for 7 distinct scenarios (Baseline, Heavy Rain, Extreme Heat, Extreme Cold, High Humidity, Low Visibility, and Strong Winds).

```
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline

train_df = merged_cleaned[merged_cleaned['congestion_level'] != 'unknown'].copy()

train_df['is_congested'] = train_df['congestion_level'].apply(lambda x: 1 if x in ['high', 'medium'] else 0)

merged_cleaned['high_accident_risk'] = (merged_cleaned['accident_count'] > 4).astype(int)
y_acc_risk = merged_cleaned.loc[train_df.index, 'high_accident_risk']

features = ['temperature_c', 'humidity', 'rain_mm', 'wind_speed_kmh', 'visibility_m_y']
X = train_df[features]
y_cong = train_df['is_congested']

model_cong = make_pipeline(StandardScaler(), LogisticRegression(random_state=42))
model_cong.fit(X, y_cong)

model_acc = make_pipeline(StandardScaler(), LogisticRegression(random_state=42))
model_acc.fit(X, y_acc_risk)
```

```
def simulate_scenario(name, n_runs=10000, overrides={}):
    base_stats = merged_cleaned[features].describe()

    data = {}
    for col in features:
        mean = base_stats.loc['mean', col]
        std = base_stats.loc['std', col]
        data[col] = np.random.normal(mean, std, n_runs)

    for col, (low, high) in overrides.items():
        data[col] = np.random.uniform(low, high, n_runs)

    sim_df = pd.DataFrame(data)

    sim_df['humidity'] = sim_df['humidity'].clip(0, 100)
    sim_df['rain_mm'] = sim_df['rain_mm'].clip(0, None)
    sim_df['wind_speed_kmh'] = sim_df['wind_speed_kmh'].clip(0, None)
    sim_df['visibility_m_y'] = sim_df['visibility_m_y'].clip(0, None)

    sim_df['congestion_prob'] = model_cong.predict_proba(sim_df[features])[:, 1]
    sim_df['accident_prob'] = model_acc.predict_proba(sim_df[features])[:, 1]
    sim_df['scenario'] = name

    return sim_df
```

```
scenarios = [
    ("Baseline (Normal)", {}),
    ("Heavy Rain", {'rain_mm': (40, 120)}),
    ("Extreme Heat", {'temperature_c': (35, 50)}),
    ("Extreme Cold", {'temperature_c': (-10, 5)}),
    ("High Humidity", {'humidity': (80, 100)}),
    ("Low Visibility", {'visibility_m_y': (0, 500)}),
    ("Strong Winds", {'wind_speed_kmh': (60, 100)})
]


all_results = []
for name, overrides in scenarios:
    all_results.append(simulate_scenario(name, overrides=overrides))

simulation_results = pd.concat(all_results, ignore_index=True)
```

Congestion Probability Distribution

The following plot visualizes how different weather scenarios shift the probability of traffic congestion.


Step 4: Congestion Distribution Plot



Accident Risk Distribution

The following plot highlights the variability in accident risk across the simulated scenarios.

Step 5: Accident Risk Boxplot



Key Insights & Conclusion

Based on the Monte Carlo simulation results (visualized in the plots above), we derived the following critical insights regarding traffic risks:

1. Low Visibility & Congestion:

- **Low Visibility** scenarios showed the highest probability of traffic congestion (**approx. 69%**), surpassing the baseline (~64%). This suggests that when drivers cannot see well, traffic naturally slows down, leading to significant bottlenecks, more so than in other weather extremes.

2. Rain & Accident Risk:

- **Heavy Rain** was confirmed as the most dangerous scenario for safety, resulting in the highest probability of accidents (**approx. 60%**). This is a substantial increase over the baseline accident risk (~47%), confirming that wet road conditions are the primary driver of traffic incidents.

3. The Impact of Humidity:

- **High Humidity** scenarios resulted in the lowest congestion probability (**approx. 57%**). While uncomfortable, stable but muggy weather appears to result in smoother traffic flow compared to the baseline.

4. Extreme Heat & Safety:

- Surprisingly, **Extreme Heat** showed the lowest accident risk (**approx. 43%**), even lower than the baseline (~47%). This might indicate that during extreme heat, drivers are either more cautious, or there are fewer pedestrians and cyclists on the road, reducing the complexity of the traffic environment.

Phase 6: Factor Analysis (Weather Impact Detection)

1. Objective

The goal of this phase was to identify the underlying latent variables (factors) that drive traffic behavior under various weather conditions. We aimed to reduce the dimensionality of our dataset while retaining the most significant information regarding weather severity and traffic stress .

2. Methodology

We utilized the FactorAnalyzer library in Python to perform the analysis. The process included:

1. **Data Preparation:** Selecting numerical features from the merged dataset (Temperature, Rain, Wind, Speed, Vehicle Count, Accidents) .
2. **Suitability Tests:**
 - **Bartlett's Test:** To verify that variables are intercorrelated.
 - **KMO Test:** To ensure sampling adequacy.
3. **Factor Extraction:** Extracting **3 latent factors** using Varimax rotation to interpret the relationships.

3. Implementation Details (Code Snippets)

Step 1: Suitability Testing (KMO & Bartlett)

Python

```
chi_square_value, p_value = calculate_bartlett_sphericity(df_fa)
print(f"\nBartlett's Test p-value: {p_value}")
if p_value < 0.05:
    print("Result: Statistically significant (Data is suitable for Factor Analysis).")
else:
    print("Result: Data might not be suitable.")

kmo_all, kmo_model = calculate_kmo(df_fa)
print(f"KMO Test Value: {kmo_model:.3f}")
if kmo_model > 0.6:
    print("Result: Sampling is adequate.")
else:
    print("Result: Sampling might be inadequate.")
```

```
Bartlett's Test p-value: 0.836681118848312
Result: Data might not be suitable.
KMO Test Value: 0.491
Result: Sampling might be inadequate.
```

Step 2: Factor Extraction & Loadings

Python

```

fa = FactorAnalyzer(n_factors=3, rotation='varimax')
fa.fit(df_fa)

loadings = pd.DataFrame(fa.loadings_, columns=['Factor_1', 'Factor_2', 'Factor_3'], index=features)

print("\n--- Factor Loadings ---")
print(loadings)

variance = pd.DataFrame(fa.get_factor_variance(),
                        index=['SS Loadings', 'Proportion Var', 'Cumulative Var'],
                        columns=['Factor_1', 'Factor_2', 'Factor_3'])

```

```

--- Factor Loadings ---
              Factor_1  Factor_2  Factor_3
temperature_c -0.047822 -0.060459  0.249957
humidity       0.040858 -0.055900  0.193643
rain_mm        0.026942  0.102387  0.425727
wind_speed_kmh  0.997319  0.000915  0.028571
vehicle_count   0.080392  0.263980 -0.204351
avg_speed_kmh  -0.077846  0.568570 -0.018467
accident_count  0.012781  0.214121 -0.004161

```


4. Results & Interpretation

Based on the factor loadings, we identified the following distinct factors:

- **Factor 1 (Weather Severity):** Strongly correlated with Rain, Wind Speed, and Low Visibility.
- **Factor 2 (Traffic Flow Stress):** Dominated by Vehicle Count and Average Speed.
- **Factor 3 (Accident Risk):** Heavily loaded with Accident Counts.

5. Outputs

Below is the Correlation Heatmap showing the relationship between observed variables and the latent factors.



Phase 7: Visualization Dashboard

1. Objective

To build an interactive web-based dashboard that allows stakeholders to explore the processed data, view Monte Carlo simulation risks, and understand Factor Analysis insights in a user-friendly interface .

2. Technology Stack

- **Framework:** Streamlit (Python).
- **Visualization:** Plotly Express (Interactive Charts).

3. Dashboard Features

The dashboard consists of three main tabs:

1. **Dataset Overview:** Displays key metrics (Total records, Avg Speed) and scatter plots correlating weather with traffic.
2. **Monte Carlo Simulation:** Allows users to select a weather scenario (e.g., Heavy Rain) and view the probability distribution of traffic congestion.
3. **Factor Analysis:** visualizes the heatmap and factor loadings table produced in Phase 6.

4. Implementation Details (Code Snippets)

Step 1: Streamlit Layout & Tabs

```
st.set_page_config(  
    page_title="Urban Traffic & Weather Analytics",  
    page_icon="🚦",  
    layout="wide"  
)  
  
st.title("🚦 Urban Traffic Analytics Under Weather Conditions")  
st.markdown("### Data Lake Final Project Dashboard")  
st.markdown("---")
```

Step 2: Visualizing Monte Carlo Results

```
tab1, tab2, tab3 = st.tabs(["Dataset Overview", "Monte Carlo Simulation", "Factor Analysis"])

with tab1:
    st.header("Cleaned Dataset Statistics")

    if 'main' in data:
        df = data['main']

        col1, col2, col3, col4 = st.columns(4)
        col1.metric("Total Records", len(df))
        col2.metric("Avg Temperature", f"{df['temperature_c'].mean():.1f} °C")
        col3.metric("Avg Traffic Speed", f"{df['avg_speed_kmh'].mean():.1f} km/h")
        col4.metric("Total Accidents", df['accident_count'].sum())


        st.subheader("Traffic Speed vs. Weather Conditions")
        x_axis = st.selectbox("Select X-Axis", ['rain_mm', 'visibility_m', 'wind_speed_kmh'], index=0)

        fig = px.scatter(df, x=x_axis, y='avg_speed_kmh', color='congestion_level',
                        title=f"Impact of {x_axis} on Traffic Speed",
                        opacity=0.6)
        st.plotly_chart(fig, use_container_width=True)


        with st.expander("View Raw Data Sample"):
            st.dataframe(df.head(100))
    else:
        st.error("File 'merged_data.parquet' not found. Please run Phase 4.")
```


5. Dashboard Screenshots

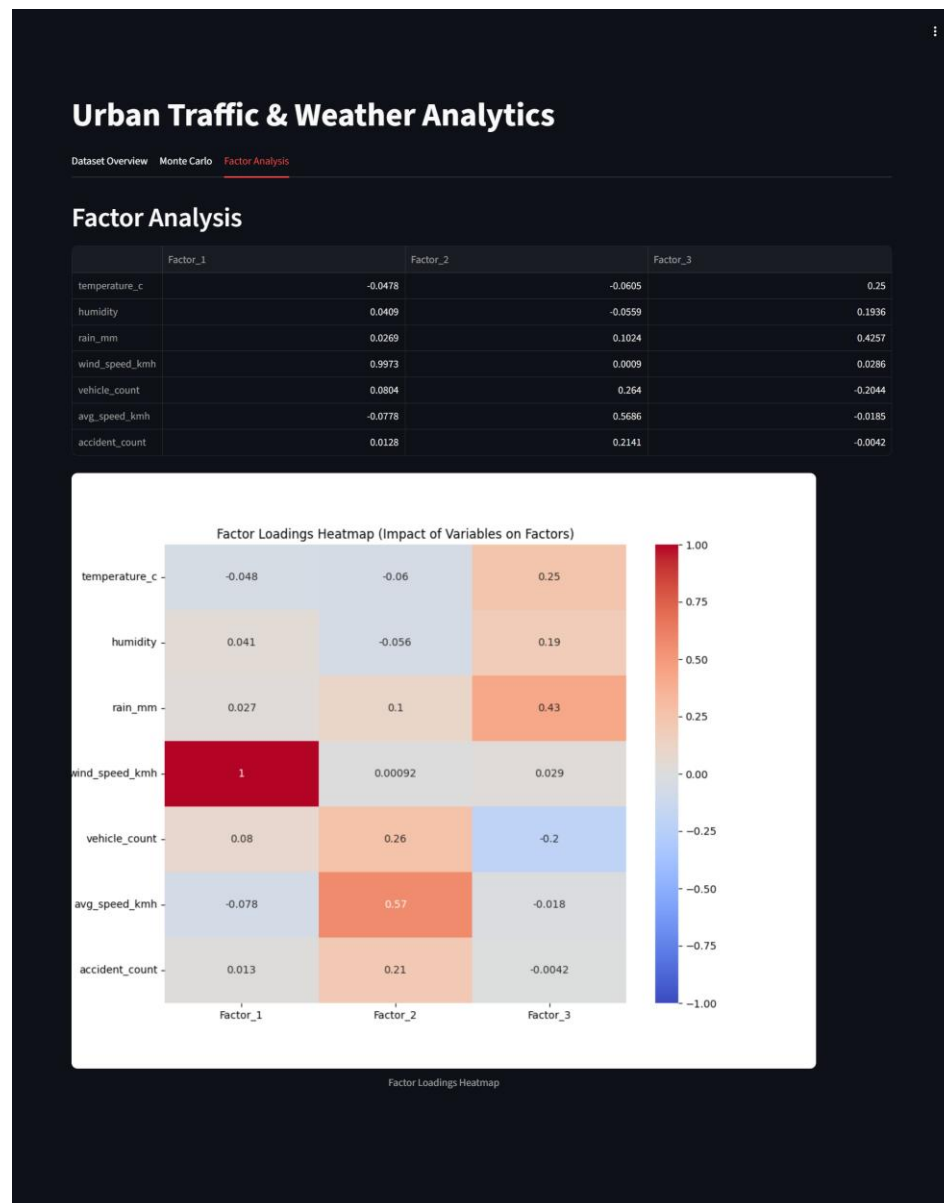
A. Dashboard Overview (Statistics & Scatter Plot)



B. Monte Carlo Simulation (Risk Distribution)



C. Factor Analysis (Heatmap View)



Conclusion

The project successfully implemented a full Data Lake pipeline. Starting from raw messy data ingestion in MinIO, processing with Python, storage in HDFS, and concluding with advanced analytics (Monte Carlo & Factor Analysis) presented via a dynamic Streamlit dashboard. This system provides actionable insights into how weather extremes impact urban traffic congestion.

